

Title

Abandoning Objective Abandoning Objectives: Evolution Through
Through the Search for Novelty Alone

Joel Lehman University of Central Florida Kenneth O. Stanley
University of Central Florida

Colin Milhaupt & Greg Jacobus University of New Mexico

Introduction

- ▶ Discusses algorithms used for evolutionary computation (EC).
- ▶ Previous approaches primarily used objective functions.
 - ▶ Example: Chinese finger trap
 - ▶ Problem: Only finds local optima (deception)
- ▶ New Idea: search for novelty instead.
 - ▶ Completely ignores the objective
- ▶ How can we use these approaches to model evolution?
- ▶ Main idea: The objective function isn't as perfect as we used to think, sometimes novelty functions work better.



Figure 1: Chinese Finger Trap

Background

- ▶ Deception – when lower-order building blocks are combined and don't lead to a global optimum.
 - ▶ Multi-Objective Evolutionary Algorithms can sometimes fight deception by training against multiple objectives.
 - ▶ Incremental evolution of objective functions can also help
 - ▶ Implies the need to track the “stepping stones” that lead to the objective
- ▶ Selection pressure restricts the scope and direction of search
- ▶ NeuroEvolution of Augmenting Topologies (NEAT)
 - ▶ Evolves artificial neural networks
 - ▶ Complexifies the network over generations

The Search for Novelty

- ▶ Learning method is rewarded finding instances that are significantly different from any found before
- ▶ Biped locomotion example:
 - ▶ Novelty function would reward falling in new ways (maybe eventually walking)
 - ▶ Objective function function would reward falling the furthest
- ▶ Novelty search is different from exhaustive search
 - ▶ Domain typically limits the variety of behaviors
 - ▶ Since NEAT starts simple and complexifies, it is much better than random.

Novelty Search Algorithm

- ▶ point

Experiment 1

- ▶ point

Biped Experiment

- ▶ Video

Discussion/Conclusion

- ▶ point